

# FAT32 Data Layout Specification

Thank you OSDev & Wikipedia

component	offset	size	description
BIOS Parameter Block	0	3	Literal: EB 3C 90 Comment: used to prevent data execution in boot environment
	3	8	Literal: MSWIN4.1 Comment: ignored, but keep it for consistency with standard
	11	2	Comment: bytes per sector Default value: 512 bytes
	13	1	Comment: Sectors per cluster Default: 64; this gives 32Kb clusters
	14	2	Comment: # of reserved sectors; presumably 2? Not sure if FAT is included
	16	1	Comment: # of FATs Default value: 2 (duplicate FAT for redundancy)
	17	2	Comment: # of FAT12 or FAT16 root directory entries Default: 0; this is FAT32
	19	2	Comment: Sectors on the volume; 0 if > 65535
	21	1	Comment: Media Descriptor Type Default: F8
	22	2	Comment: Sectors Per FAT (for FAT16 and 12) Default value: 0; always ignored
	24	2	Physical sectors per track Default value: 0; always ignored
	26	2	Number of heads per disk Default value: 0
	28	4	Count of hidden sectors preceding volume; used when partitioning is used Default value: 0; we're not partitioning
	32	4	Comment: large amount of sectors on media; used if sector count > 65535 See offset 19
FAT32 extended fields	36	4	Clusters per FAT
	40	2	Flags Default: 0; ignored

	42	2	FAT version Default: ???
	44	4	Cluster # of root directory Default: 2, because OSDev says so
	48	2	Sector # for FSInfo structure
	50	2	Sector # for backup boot sector
	52	12	Reserved Default: 0
	64	1	Driver number Default: 80 (generally ignored otherwise)
	65	1	Flags for Windows NT Default: 0; ignored generally
	66	1	Signature Default: 0x28
	67	4	VolumeID serial # Default: 0; generally ignored
	71	11	Volume Label String; padded with spaces (0x20) Default: "AlphaGoS"
	82	8	System identifier string Default: "FAT32 "; ignore this
	90	420	Boot code; default: empty
	510	2	Bootable partition signature. Default: 0xAA55
FS INFO Sector	0 (starting in sector 1)	4	Compatibility guard for earlier versions of FAT Default: 0x52 0x52 0x61 0x41
	4	480	Reserved; default to 0
	484	4	Info sector signature; Default 0x72 0x72 0x41 0x61
	488	4	Last known number of free data clusters; should not rely on this; defaults to 0xFFFFFFFF; should be ignored for this implementation
	492	4	Number of most recently known allocated data cluster; set to 0xFFFFFFFF. Ignore
	496	12	Reserved; set to 0
	508	4	Sector signature. Set to 0x00 0x00 0x55 0xAA

## FAT Table

Entries are 32 bit values. #of entries is determined by value at offset 36 of above. Note that there are two FATs back to back.

Value	Meaning	Comment
0x?0000000	Free Cluster	
0x?0000001	Reserved Cluster	
0x?0000002 - 0x?FFFFFFEF	Used cluster; value points to next cluster	
0x?FFFFFFF0 - 0x?FFFFFFF6	Reserved	
0x?FFFFFFF7	Bad Cluster (do not use)	
0x?FFFFFFF8-0x?FFFFFFF	Last cluster in file	Once this is hit, stop reading the file

The first two entries in the FAT are special:

- Cluster zero: 0x011111F0
- Cluster 1: 0x0FFFFFFF

Data begins immediately after last FAT table

Per specification above (see offset 44) the root directory is at cluster 2

## Directory Entries

32 byte entries

1024 entries per cluster

Just Data

Offset	Length	Comment	Values Details
0	11	8.3 file name; first 8 are name; last 3 are ext	First byte 0 ->no more entries First byte 0xE5 -> unused entry First byte 0x1 -> used entry All other

			values: undefined; field unused for compliance
11	1	Attributes; 0x01: readonly; 0x02: hidden; 0x04: system; 0x08: volume_id Above flags indicate this is associated with a long-file-name entry 0x10: directory; 0x20=archive	
12	1	Reserved for use by Windows NT	
13	1	Creation time in tenths of a second; 0-199 inclusive	
14	2	Creation time: Hour: 5bits, Minutes:6bits: seconds: 5bits (multiply by 2)	
16	2	Creation date: Year:7bits;Month:4bits;Day:5bits	
18	2	Last accessed date. See format above	
20	2	High 16 bits of cluster number	
22	2	Last modification time	
24	2	Last modification date	
26	2	Low 16 bits of cluster number	
28	4	Size of file in bytes	

### Long file names:

- Long file names always precede the short-file-name version
- Any number of long file name entries may exist for a file or directory
- Because multiple name blocks may exist, they contain an ordering field; they are not guaranteed to appear in correct display order on disk

Offset	Length	Comment
0	1	Order (index) in the overall long file name

1	10	First 5 2-byte (think unicode) characters
11	1	Always 0x0F (used to identify that this is a name block)
12	1	Zero (for name entry)
13	1	Checksum of the short file name; ignore
14	12	Next 6 2-byte characters
26	2	Always Zero
28	4	Final 2 2-byte characters of the name

## FAT32/VFAT Patent Compliance

The process of using directory entries to store extended names in addition to the 8.3 filenames is patented. To avoid infringing on the 100% completely valid and sensible patent, this implementation will not consider the 8.3 filename

# Implementations for Common Processes on FAT32 FS

## Lookup File By Name - Long File Names

Assumes starting from root

- Go to root directory at cluster 2
- Iterate through 32-byte entries
  - If this is a long-name entry
    - Read up to short name entry
    - Assemble name; don't forget to account for 2-byte vs 1-byte chars
    - Compare both long and short names
  - If end of cluster reached AND cluster is marked as having a next cluster
    - Jump to that cluster
    - Continue search
    - Note: you could jump clusters while assembling a long-file-name

## Calculating Raw-Byte offsets for clusters

- Simple driver specification says that data is accessed by byte offset from 0
- Note: on 32-bit systems, this limits us to 4Gb

Notes:

- Sectors are 512 bytes; not apparently relevant here
- FAT Table is offset by 2 sectors

- Clusters are 32Kb each
- FAT table consumes 32bits \* Cluster-per-fat space
- Note that there are 2 FATs (most likely)
- Clusters 0 and 1 are reserved; data starts at cluster 2

Byte\_for\_cluster = 2\*sector\_size + 4\*cluster\_per\_fat\*2 + 32768\*cluster  
 = 1024 + 8\*cluster\_count + 32768\*cluster\_number

## Calculating Raw-Byte offsets to get to the FAT

- FAT0: 1024
- FAT1: 1024 + sizeof(FAT0)

Size of a FAT is 4\*clusters\_per\_fat

## Implementing Query on top of the Simple Driver Specification

Supported metadata properties:

- Created
- Last updated
- Is Directory
- Readonly
- Hidden
- System
- Last read
- Size in bytes

Side effects

- Updates last read

## Implementing Set on top of the Simple Driver Specification

Supported metadata properties:

- Created - not settable
- Last updated - not settable
- Is Directory - not settable
- Readonly
- Hidden
- System
- Last read - not settable
- Size in bytes - not settable
- Name

Side Effects

- Updates last updated

Cautions:

- Changing the name may have dramatic implications for directory table in case of long file names

General algorithm:

1. Recursively read filesystem until directory entry for file is found
2. If it is a regular property, update it.
3. Else if the name is being updated:
  - a. Determine # of entries currently being used for name
  - b. If # of entries needed for new name is same as old name, update in place
  - c. If the # of entries differs, construct new entries in memory
    - i. Then shift table up
    - ii. Append entry at end of directory
    - iii. Note: may need to jump to next cluster for any part of this operation
4. Update last-updated field

## Implementing Read on top of the Simple Driver Specification

Cautions:

- Be careful about managing # of bytes read when jumping between clusters

Side Effects

- Updates last read

General algorithm:

1. Recursively read filesystem until directory entry for file is found
2. Read the cluster and the size of the file
  - a. Use the file offset to determine start point in file and jump clusters as needed
  - b. Read into a buffer until end of cluster is reached, then jump
  - c. Continue reading until bytes read count is consistent with request or end of file is reached  
(offset + bytes\_read == file\_size)

Note: reading is finished when the current position is cluster\_start

## Implementing Create on top of the Simple Driver Specification

Cautions:

- Don't forget to update duplicate FAT

Side Effects

- Creates entry in directory table
- Updates FAT table

General Algorithm:

1. Recursively find the directory entry to which this will be added
2. Construct long file name entry
3. Construct short file name entry
4. Grab a free cluster from FAT, update entries

# Implementing Delete on top of the Simple Driver Specification

Cautions:

- Make sure to delete associated long file name entries
- Don't delete the root directory
- When updating FAT, don't forget to update duplicate FAT

Side effects:

- Removes entries (potentially many) from FAT table
- Removes entries (potentially many) from Directory entry

General Algorithm:

1. Recursively find the directory entry for this
2. Recursively (accounting for jumps) free associated cluster entries in the FAT
3. Mark directory entry as unused

# Implementing Write on top of the Simple Driver Specification

Cautions:

- When updating FAT, don't forget to update duplicate FAT

Side effects:

- Updates last updated in directory entry
- Updates size in bytes in directory table
- May require pulling more clusters from the FAT table and updating the existing ones

General Algorithm:

1. Recursively read filesystem until directory entry for file is found
2. Read the cluster and the size of the file
  - a. Use the file offset to determine start point in file and jump clusters as needed
  - b. Write from buffer until end of cluster is reached or end of buffer is reached
    - i. If end of cluster is reached, grab a new cluster, update FATs, continue writing at next FAT

# Implementing the Formatter on top of the RAW File System

Because an unformatted device does not have the information needed by FAT, it must first be formatted using a separate utility. To do this from within the OS:

- Mount the device with RawFS
- Run the formatter utility on the mounted directory
- Unmount the directory
- Test and mount with FAT32FS



Issues: need to decide on number of clusters based on size of device

- Fixed overhead: 1024 bytes (1Kb)
- Cluster size: 32kbytes
- Cluster count:  $((\text{device\_size\_kbytes} - 1\text{Kb}) / 32\text{kbytes}) + 2 - 1$ 
  - -1 above accounts for size of FAT table ( it is at most 4kb)



